

Accessing Databases with R

Sean Hellingman ©

Data Visualization and Manipulation through Scripting (ADSC1010)

shellingman@tru.ca

Fall 2024



THOMPSON RIVERS UNIVERSITY

Topics

- 2 Introduction
- 3 Relational Databases
- 4 Accessing a Database from R
- 5 Working With the Data
- 6 Joining Data
- 7 Process
- 8 Exercises and References

Introduction

- There are R packages that allows users to connect to and query a database directly.
- This allows you to use R syntax and data structures to work with databases.
- We can use the *dbplyr* package to help with this.

Databases

- Simple data sets are easily saved and accessed as .csv files.
- As the complexity of data grows, you may need multiple frames to organise your data.
- Difficult to store data of different structures in a single sheet.
- The data files may simply be too large to store on a local computer.
- A database (database management system) is used to organise, save, and access information.

Relational Databases

- Relational databases are the most commonly used.
- Data is organised into tables where each row represents a **record** and each column represents a **field** (individual data property of that item).
- Tables are structured like data frames in R.
- Databases may have thousands of tables representing different facets of the data.
- **Relational database management system (RDMS)**

Primary Key

- Relational databases identifies each record in the database table using a **primary key**.
- In each table, one field (column) is designated as the primary key that is unique to each row.
- Primary keys are unique *identifiers* for each observation in the data.
- As data within a database can change, we cannot use row numbers as a primary key.
- Only one primary key is permitted per table.

Foreign Key

- Each *record* may be associated with another.
- Example: Assume we have table with information about musical artists and another about individual songs.
 - We can connect the two tables based on the names of the musical artists.
- **Foreign keys** allow you to join tables together like you would using the `join()` function in R.
- Provides the *relational* functionality of relational databases.

Comments

- There are many different Structured Query Language (SQL) developers.
- SQLite is the simplest database system, generally not used in industrial settings.
- Others free developers: PostgreSQL & MySQL.

Motivation

- You can access your data directly from the database, query the data you want, save that data, and then import it into R or some other statistical software.
- Or, we can use R to directly query a database directly.
- Then we can use familiar R syntax to work with databases.
- More specifically, we will use the functionality of the *dplyr* package to manipulate the data.

R Packages

- Need to use *dplyr*.
- Will need to install and load the *dbplyr* package.
- Will need to install and load the *DBI* package.
 - *The DBI package helps connecting R to database management systems.*
- Install and load the *RSQLite* package.
 - *If you want to access SQLite databases*
- Install and load the *RPostgreSQL* package.
 - *If you want to access Postgres databases*
- Remeber: `if(!require(package_name))
install.packages("package_name") library(package_name)`

Connecting to Databases

- Databases are managed and accessed through an RDMS, which is separate from R.
- We need to “connect” to the database through R.
- Use: `db_connection <- dbConnect(SQLite(), dbname = "path/to/database.sqlite")`
 - First argument is the relevant database connection package (RSQLite).
 - If database is in working directory, you can just access the database directly.
 - User & password: `user="Username", password="Password123"`
- To disconnect from your database use:
`dbDisconnect(db_connection)`

Example 1

- Download the *Chinook_Sqlite.sqlite* database from moodle and save it in an appropriate location.
- Use the `dbConnect()` function to access the *Chinook_Sqlite.sqlite* database from R.

Chinook Database

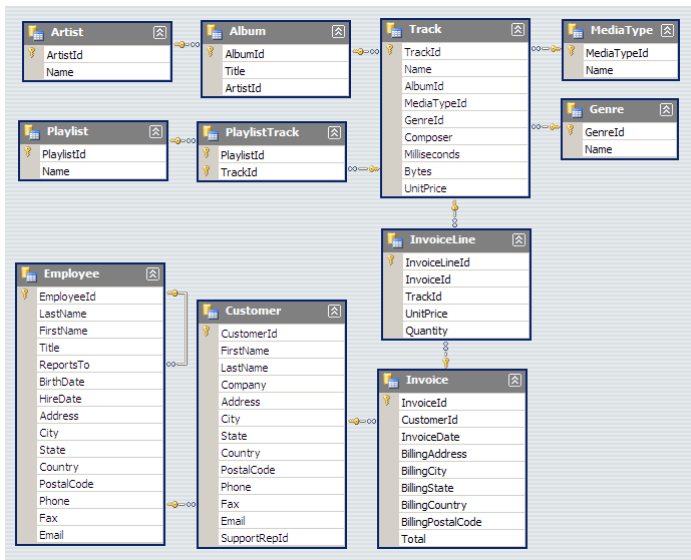


Table Names

- Once you have accessed the database, you can use the `dbListTables()` function to get a vector of the table names.
- Remember: data come from specific tables within the database.
- Need to create a variable in R that references specific tables.
 - Can use the `tbl()` function to accomplish this.
- If you examine the variable name, it looks *mostly* like a normal data frame.
- But this variable actually comes from a remote source.

Example 2

- Use the `dbListTables()` functions to get a list of tables in the *Chinook* database.
 - You may refer to slide 13 to help you visualise the structure of the database.
- Use the `tbl()` function to create a variable for the *Track* table in R.
- What does the variable look like?

Using *dplyr*

- Once we have created a reference to the table in R, we can apply the *dplyr* functions!
- We can also construct a query using *dplyr* and generate the corresponding SQL query.
 - Use the `show_query()` function.
- Save the query as an object in R and then use the `show_query()` function.

Example 3

- Return the track *Name* and *TrackId* for all U2 Tracks (*Composer*) found in the Track table.
- Generate the corresponding SQL query.
- Use the `data.frame()` to create a data frame of U2 tracks.

left_join()

- Looks for matching columns between two data frames (tables).
- Returns a new data frame that is the first (*left*) argument with extra columns from the second (*right*) added on.
- The resulting table is a *merged* table of the two arguments.
- Matching occurs using the `by` argument which takes a vector of column names (strings).
 - Use: `by = join_by("key.1" == "key.2")` if foreign keys have different names.
- Left rows without a match will have NA in the right columns.

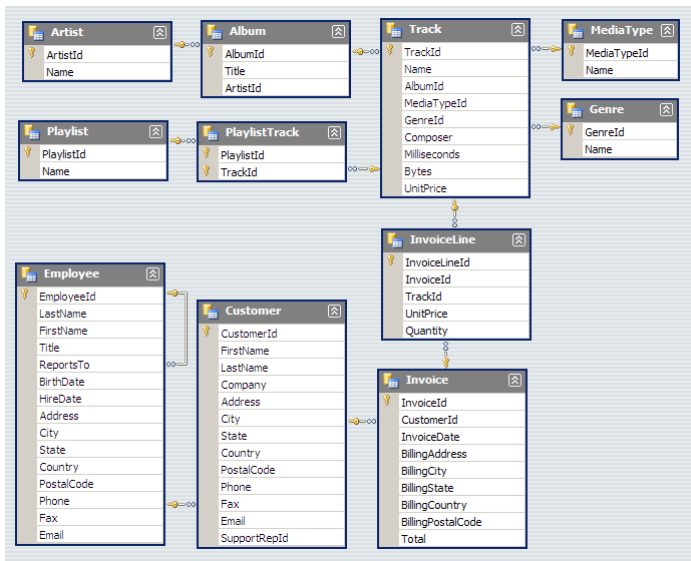
inner_join()

- Only rows present in *both* data frames (tables) are returned.
- Returns a new data frame that contains only observations that had matches in both data frames.
- Observations without matches will not be included (no NA values).
- The order of the arguments does not matter.

full_join()

- All the rows present in *both* data frames (tables) are returned.
- A row for every single observation is returned.
- Observations without matches will have NA values in the columns from the other data frame.
- Can lead to very messy data.
- The order of the arguments does not matter.

Chinook Database Again



Example 4

- Use the `tbl()` function to create a variable for the *Album* table in R.
- Include the album information to each of the tracks in the *Track* table (join).
- Use the `data.frame()` to create a data frame of your resulting table.

Process

- ❶ Create a connection to an RDMS (SQLite):
 - `db_connection <- dbConnect(SQLite(), dbname = "path/to/database.sqlite")`
- ❷ Access a specific table within the database:
 - `some_table <- tbl(db_connection, "TABLE_NAME")`
- ❸ Construct a query of the table using *dplyr* syntax:
 - `db_query <- some_table %>% filter(some_column == some_value)`
- ❹ Execute your query to return the data (bring it into R):
 - `results <- collect(db_query)`
OR
 - `results <- data.frame(db_query)`
- ❺ Disconnect from the database when you are finished:
 - `dbDisconnect(db_connection)`

Exercise 1

- Use the `dbConnect()` function to access the *Chinook_Sqlite.sqlite* database from R.
- Query all customers whose *Country* is Canada.
- Join the Canadian based customers with their respective invoices (Invoice).
 - You could have multiple rows for the same customer, if they made more than one purchase.
- Join the InvoiceLine for all of the Canadian based customers.
- *You can choose which order you would like to join the tables.*
- Import your data into R and disconnect from the database.

Exercise 2

- Using the Chinook database:
 - What is the title of the album with AlbumId 45?
 - Find the name and length (in seconds) of all tracks that have length between 55 and 75 seconds.
 - Provide a query showing a unique/distinct list of billing countries from the Invoice table.
 - Find the CustomerID of the customer(s) who made the most purchases.

References & Resources

- ① Michael Freeman, Joel Ross, *Programming Skills for Data Science: Start Writing Code to Wrangle, Analyze, and Visualize Data with R*, 2019, ISBN-13: 978-0-13-513310-1
- <https://dbi.r-dbi.org/>
- <https://github.com/lerocha/chinook-database>
- <https://dplyr.tidyverse.org/>