

# Data Manipulation in R with *dplyr*

---

Sean Hellingman ©

Data Visualization and Manipulation through Scripting (ADSC1010)

*shellingman@tru.ca*

Fall 2024



**THOMPSON RIVERS UNIVERSITY**

# Topics

- 2 Introduction
- 3 The select() Function
- 4 The filter() Function
- 5 The mutate() Function
- 6 The arrange() Function
- 7 The summarize() Function
- 8 Sequential Operations
- 9 The group\_by() Function

# Introduction

- The *dplyr* package is the primary data wrangling tool in R.
- Intended to have intuitive vocabulary.
- Once mastered, it will make your data processing much faster.

## *dplyr* Functions (Verbs)

- The package is written in a way that you should be able to use specific functions that match your task description:
  - **Select** specific columns (features) from a dataset.
  - **Filter** out any unwanted observations (rows) from your dataset.
  - **Mutate** a dataset by adding more columns (features).
  - **Arrange** observations (rows) in a particular order.
  - **Summarize** data in aggregates such as *mean*, *median*, or *minimum*.
  - **Join** multiple datasets together into a single data frame.

## Package Information

- The *dplyr* package is part of the *tidyverse* collection of R packages.
- We need to install and load the package prior to use:
  - `if(!require(dplyr)) install.packages("dplyr")`  
`library(dplyr)`
- Once the package is loaded we can use the existing functions as we please.

## select()

- The `select()` function takes the data frame and which columns you want to select as arguments.
  - `select(df.name, variable1.name, variable2.name,...)`
  - `select(df.name, - variable1.name, - variable2.name,...)`  
*selects all columns apart from those indicated with -*
- Less involved than using base R.
- Returns a data frame even if you only ask for one column.

## Example 1

- Import the the *starwars* data from the *dplyr* package.
- Use the `select()` function to create a data frame with the *name*, *height*, and *homeworld* of characters from the *starwars* data.
- Use the `select()` function to create a data frame **without** the *vehicles*, and *starships* of characters from the *starwars* data.

## filter()

- The `filter()` function takes the data frame and allows you to extract specific rows.
  - `filter(df.name, variable1.name == Value)`
  - `filter(df.name, variable2.name <= Value)`
- Takes a list of conditions and returns a data frame.



## Example 2

- Use the `filter()` function to create a data frame of all the humans from the *species* variable in the *starwars* data from the *dplyr* package.

## mutate()

- The `mutate()` function takes the data frame and allows you to create new columns.
  - `mutate(df.name, New.variable.name = Value)`
  - `mutate(df.name, New.variable.name = variable1.name + variable2.name)`
- Takes a list of new columns to create and returns a data frame.
- Note: the function does not actually add the new columns. Instead you need to replace your old data frame with the *new* one.

## Example 3

- Use the `mutate()` function to add the Body Mass Index (BMI) of the characters to the *starwars* data from the *dplyr* package.
- BMI:  $\text{kg/m}^2$

## arrange()

- The `arrange()` function takes the data frame and allows you to sort the rows of your data by some column.
  - `arrange(df.name, -variable2.name)` decreasing order
  - `arrange(df.name, variable1.name)` increasing order
- Takes a list of columns to arrange by and returns a data frame.
- The order of your list of columns matters.
- Note: you again need to replace your old data frame with the *new* one.

## Example 4

- Use the `arrange()` function to arrange the characters from the *starwars* data from the *dplyr* package by *height* in decreasing order.

## summarize()

- The `summarize()/summarise()` function acts as an aggregation operation for one or more columns from your data frame.
- It reduces the entire column into a single value.
  - `summarize(df.name, Var.2_mean = mean(variable2.name))`  
mean of *variable2.name*
  - `summarize(df.name, Var.1_max = max(variable1.name))`  
maximum value of *variable1.name*
- Takes a list of arguments and returns a data frame with one row.
- You can also write your own summary functions and use them within the `summarize()` function.
- Not usually ideal to replace your existing data with these results.

## Example 5

- Use the `summarize()` function to obtain the maximum, minimum, average, and median BMI of the characters from the *starwars* data from the *dplyr* package by *height* in decreasing order.
- You will need to omit the missing observations.

## Sequential Operations

- In order to conduct more complex analysis you may need to combine some of these functions.
- One approach is to create *intermediary* variables and pass them from one function to another.
- A more efficient approach is to use *anonymous* variables and nest the statements within other functions.



## Example 6

- Use the functions in the *dplyr* package to determine the maximum BMI of the humans with blue eyes in the *starwars* data.
  - Use *intermediary* variables.
  - Use *anonymous* variables and nest the statements.

## The Pipe Operator %>%

- Built in *dplyr* functionality that makes passing results of functions easier.
- Formally, the pipe operator %>% takes the result from one function and passes as the first argument to the next function.
- Works well as all of the *dplyr* functions take a data frame as the first argument.
- In R Studio: ctrl+shift+m results in %>%

## Example 7

- Use the pipe operator `%>%` and functions in the *dplyr* package to determine the maximum BMI of the humans with blue eyes in the *starwars* data.
- Which approach seemed easiest?

## group\_by()

- The `summarize()` function works on all observations in a column.
- The `group_by()` function allows you to group data with the same variable value together.
- Returns a **tibble** which is a special version of a data frame used primarily in the *tidyverse* packages.
- Functions (*verbs*) applied directly to the **tibble** will be applied to each group separately.

## Example 7

- Use the `group_by()` function to summarize *height* and *mass* by *eye\_color* groups in the *starwars* data.

## Data in Multiple Files

- Data is often stored in multiple multiple locations.
- Very typical of relational databases.
- Often more efficient to store and update data in this format.
- We will need to **join** data frames together to work with them.

## Joining Data Frames

- Very similar approach to SQL.
- Use columns that are present in both data frames to *match* corresponding rows together.
- Columns used in the matching process are called **identifiers** (keys).
- The identifiers are combined in one row when the data frames are joined.

## left\_join()

- Looks for matching columns between two data frames.
- Returns a new data frame that is the first (*left*) argument with extra columns from the second ("right") added on.
- The resulting table is a *merged* table of the two arguments.
- Matching occurs using the `by` argument which takes a vector of column names (strings).
- Left rows without a match will have NA in the right columns.



## `right_join()`

- Looks for matching columns between two data frames.
- Returns a new data frame that is generated in the opposite direction of the `left_join()` function.
- In other words, simply reversing the arguments from the `left_join()` function.
- *You really only need to remember how to use `left_join()` OR `right_join()`.*

## inner\_join()

- Only rows present in *both* data frames are returned.
- Returns a new data frame that contains only observations that had matches in both data frames.
- Observations without matches will not be included (no NA values).
- The order of the arguments does not matter.

## `full_join()`

- All the rows present in *both* data frames are returned.
- A row for every single observation is returned.
- Observations without matches will have NA values in the columns from the other data frame.
- Can lead to very messy data.
- The order of the arguments does not matter.

# Joins

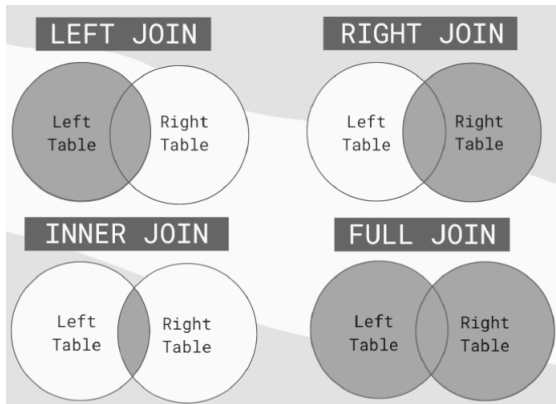


Figure: *source:* (4)

## Example 7

- Import the `Clients.csv` and `Purchases.csv` data into your workspace.
- Take some time to look through the data.
- Practice joining the data using the functions in the *dplyr* package.

## Exercise 1

- Import the *mpg* dataset from the *ggplot2* package.
- Take some time to get to know the data.
- Using the functionality of the *dplyr* package conduct the following analysis:
  - Identify the average highway miles per gallon (*hwy*) of ford vehicles newer than 1999.
  - Find the average city (*cty*) and highway miles per gallon (*hwy*) for each car *class*.
  - Create a data frame ordering the highway miles per gallon (*hwy*) from greatest to smallest for each *manufacturer*.

## Exercise 2

- Create a data frame (in R or Excel) with overlapping client IDs from Exercise 7.
- In the new data frame (randomly) include:
  - The clients' monthly incomes.
  - The clients' preferred payment method (debit or credit).
- Join your new data to the existing data and utilise the functionality in the *dplyr* package to appropriately summarise the information contained in the clients list.

## References & Resources

- ① Douglas, A., Roos, D., Mancini, F., Couto, A., & Lusseau, D. (2023). *An introduction to R*. Retrieved from <https://intro2r.com/>
- ② Michael Freeman, Joel Ross, *Programming Skills for Data Science: Start Writing Code to Wrangle, Analyze, and Visualize Data with R*, 2019, ISBN-13: 978-0-13-513310-1
- ③ Baumer, Kaplan, Horton. *Modern Data Science with R*. Retrieved from <https://mdsr-book.github.io/mdsr3e/>.
- ④ <https://learnsql.com/blog/learn-and-practice-sql-joins/>
- <https://www.tidyverse.org/>
- <https://tibble.tidyverse.org/>
- <https://statisticsglobe.com/r-dplyr-join-inner-left-right-full-semi-anti>